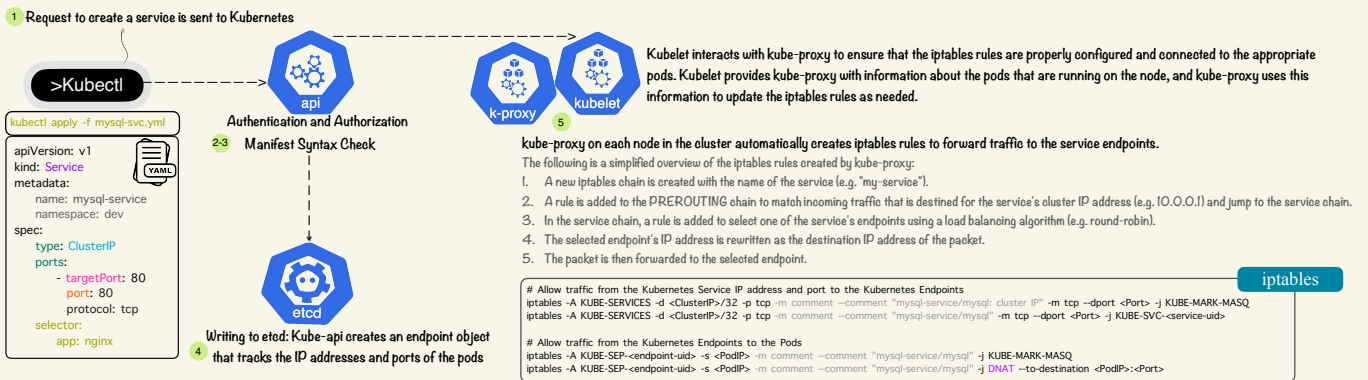


Process of creating a service in Kubernetes:

- 1 Define the service:** The first step in creating a service is to define the service using a YAML file or through the Kubernetes API. The YAML file specifies details such as the name of the service, the selector used to identify the pods that the service should route traffic to, and the type of service (ClusterIP, NodePort, or LoadBalancer).
- 2 Submit the service definition:** Once the service definition is created, it can be submitted to the Kubernetes API server
- 3 API Server validates the service definition:** The Kubernetes API server receives the service definition and validates it to ensure that it is well-formed and contains all the required information.
- 4 Service is created:** Once the service is created, Kubernetes creates an endpoint object that tracks the IP addresses and ports of the pods that the service should route traffic to. This information is stored in etcd
- 5 iptables rules are created:** Once the endpoint object is created, kube-proxy creates iptables rules on each node in the cluster to route traffic to the pods that are part of the service. These iptables rules are used to ensure that traffic is routed to the correct pod, and that traffic is load balanced across multiple pods if more than one pod matches the selector.
- 6 Access the service:** The service is now accessible within the cluster using its name or DNS name, and can be used to route traffic to the pods that are part of the service.
- 7 Monitor the service:** Once the service is running, Kubernetes monitors its health and takes action if any issues arise. For example, if a pod fails, Kubernetes will automatically remove it from the list of endpoints for the service.



Endpoint

When a Service is created, the Service controller queries the Kubernetes API server to get a list of all Pods that match the Service's label selector. It then creates an Endpoint object that includes the IP addresses and ports of these Pods, and associates the Endpoint with the Service. The Kubernetes networking layer uses this Endpoint information to route traffic to the appropriate Pods that make up the Service.

DNS

Kubernetes has a built-in DNS component that provides naming and discovery between pods running on the cluster. It assigns DNS records (A records, SRV records, etc) for each pod/service automatically. The DNS name follows a specific format, such as `<service-name>.<namespace>.svc.cluster.local` for accessing a Service or `<pod-name>.<service-name>.<namespace>.svc.cluster.local` for accessing a specific Pod associated with a headless Service.

If a Pod located in the "default" namespace needs to communicate with a service named "nginx-service" residing in the "dev" namespace, it can do so by using the URL `http://nginx-service.dev.svc.cluster.local`.

`<pod-name>.<service-name>.<namespace>.svc.cluster.local`

In Kubernetes, FQDN stands for Fully Qualified Domain Name. It is a complete domain name that specifies the exact location of a resource within the DNS hierarchy. By using FQDNs, Kubernetes simplifies the process of resource discovery, network routing, and namespace isolation within the cluster

The default DNS provider in Kubernetes is CoreDNS, which runs as pods/containers inside the cluster. CoreDNS retrieves pod/service information from the Kubernetes API to update its DNS records.

Notice: Kubernetes does not automatically create DNS records for Pod names directly. This is because Pod IPs keep changing whenever Pods are recreated or rescheduled. Instead, stable DNS records are maintained at the Service level in Kubernetes. Services have unchanging virtual IPs that act as stable endpoints

Pod dns policy

Pod's DNS settings can be configured based on the `dnsPolicy` field in a Pod specification. This `dnsPolicy` field accepts three possible values:

- ClusterFirst:** Any DNS query that does not match the configured search domains for the Pod are forwarded to the upstream nameserver. This is the default policy if `dnsPolicy` is not specified.
- None:** Allows a Pod to ignore DNS settings from the Kubernetes environment. All DNS settings are supposed to be provided using the `dnsConfig` field in the Pod Spec.
- Default:** Use the DNS settings of the node that the Pod is running on. This means it will use the same DNS as the node that the Pod runs on.

Please note that the Pod's DNS config allows you to customize the DNS parameters of a Pod

In this example, the Pod `mypod` uses a custom DNS resolver (1.2.3.4) and a custom search list (`ns1.svc.cluster-domain.example` and `my.dns.search.suffix`). The option `ndots:2` means that if a DNS query name contains less than 3 dots, then the search list mechanism will be used. For example, a query for `mypod` will be first tried as `mypod.ns1.svc.cluster-domain.example` and if that fails, as `mypod.my.dns.search.suffix`.

```
apiVersion: v1
kind: Pod
metadata:
  name: mypod
spec:
  containers:
    - name: mypod
      image: myimage
  dnsPolicy: "None"
  dnsConfig:
    nameservers:
      - 1.2.3.4
    searches:
      - ns1.svc.cluster-domain.example
      - my.dns.search.suffix
    options:
      - name: ndots
        value: "2"
```

