


**Sidecar container** is a container that is deployed alongside a main container in a pod. The main container is typically an application that performs some specific function, while the sidecar container provides support or complementary functionality to the main container.

The idea behind the sidecar pattern is to keep the main container focused on a specific task or functionality, while delegating other tasks to the sidecar container. This allows for more modular and flexible deployment architectures, as the sidecar container can be updated or replaced independently of the main container.

Although containers inside a pod share a common network and storage, they have independent lifecycles and can be created, updated, and deleted individually.




#### important use cases

- **Logging and Monitoring:** A side container can be used to collect and forward logs and metrics from the main application container to a central monitoring system.
- **Backup and Recovery:** A side container can be used to perform backup and recovery operations on the main application container.
- **Service mesh:** A sidecar container can be used to implement a service mesh such as Istio or Linkerd. A service mesh provides additional functionality for managing and securing communications between services running in Kubernetes.

One example of how a sidecar container can be used with a database service in a Kubernetes deployment:

The main container is running a database service and is exposing port 5432 for incoming database connections.

The sidecar container is configured to perform backups of the database.



The sidecar container can periodically backup the database to a remote location to ensure data resiliency.

The sidecar container is running a script that periodically backs up the database and stores the backup files in the "/backups" directory.

The script is also using the "pg\_dump" command to perform the backup and gzip to compress the backup file. The backup location is specified in the environment variable "BACKUP\_LOCATION", which is set to an S3 bucket. The script is running in an infinite loop and sleeps for 2.4 hours between each backup.

The two containers are communicating using shared volumes and environment variables. The main container is using a volume mount called "db-data" to store its data files, while the sidecar container is using a volume mount called "backup-data" to store its backup files.

```
apiVersion: v1
kind: Pod
metadata:
  name: db-pod
spec:
  containers:
    - name: db-container
      image: my-database-image
      env:
        - name: DATABASE_URL
          value: "postgresql://my-database-hostname:5432/my-database"
      ports:
        - containerPort: 5432
      volumeMounts:
        - name: db-data
          mountPath: /var/lib/postgresql/data

    - name: sidecar-container
      image: my-sidecar-image
      env:
        - name: BACKUP_LOCATION
          value: "s3://my-bucket/my-backups"
        - name: DATABASE_PASSWORD
          valueFrom:
            secretKeyRef:
              name: db-secrets
              key: database-password
      volumeMounts:
        - name: backup-data
          mountPath: /backups
        - name: db-secrets
          mountPath: /secrets
      command: ["bin/sh", "-c"]
      args:
        - |
          while true; do
            pg_dump -U postgres -h localhost my-database | gzip > /backups/my-database-$(date +%Y-%m-%d-%H%M%S).sql.gz; s3cmd put /backups/my-database-*.sql.gz "$BACKUP_LOCATION";
            sleep 86400;
            s3cmd put /backups/my-database-*.sql.gz "$BACKUP_LOCATION";
          done
      volumes:
        - name: db-data
          emptyDir: {}
        - name: backup-data
          emptyDir: {}
        - name: db-secrets
          secret:
            secretName: db-secrets
```

#### Job & CronJobs

**Job** is a type of resource that allows you to create and manage a finite or batch process in your cluster. Jobs are commonly used for tasks that need to be run once or a few times, such as **data processing**, **backups**, or **migrations**.


A Job creates one or more Pods and will continue to retry execution of the Pods until a specified number of them successfully terminate. As pods successfully complete, the Job tracks the successful completions. When a specified number of successful completions is reached, the task (ie, Job) is complete.

The backoffLimit specifies the number of times k8s should retry the Job if it fails before giving up.

```
apiVersion: batch/v1
kind: Job
metadata:
  name: data-processing-job
spec:
  backoffLimit: 3
  template:
    spec:
      containers:
        - name: data-processor
          image: data-processor:v1.4
          command: ["python", "process_data.py"]
          restartPolicy: Never
```

**CronJobs** in Kubernetes are a way to schedule and automate the execution of Jobs on a recurring basis. A Job is a Kubernetes object that creates one or more Pods to perform a specific task, and a CronJob is a higher-level abstraction that allows Jobs to be scheduled according to a specific time or interval, similar to the Unix cron utility.

Let's say you have a web application that periodically needs to generate reports based on user data. You could create a CronJob that runs a script to generate the report and then terminates when the report is complete.



CronJobs create Jobs which in turn create Pods to run the task.

Job will run at the top of every hour.

This will delete the Pod 100 seconds after it finishes.

**Notice:** By default, completed Jobs and Pods are retained after running. To automatically clean up completed Jobs, you can set ".spec.successfulJobsHistoryLimit" and ".spec.failedJobsHistoryLimit" on the CronJob.

```
apiVersion: batch/v1
kind: CronJob
metadata:
  name: report-generation-cronjob
spec:
  schedule: "0 * * * *"
  jobTemplate:
    spec:
      template:
        spec:
          ttlSecondsAfterFinished: 100
          containers:
            - name: report-generator
              image: my-django-app:v1
              env:
                - name: DJANGO_SETTINGS_MODULE
                  value: myapp.settings
              command: ["python", "manage.py", "generate_report"]
              restartPolicy: Never
```