

API Groups

In k8s, API groups are a way of organizing related resources and operations together. This allows for easier discovery and usage, and also helps to avoid naming conflicts between different resources. When k8s was first introduced, all the resources like Pod, Service, ReplicationController, etc., were all part of a single group, the "core" group, and were accessed at the path `/api/v1`. As k8s evolved and more resources were added, it became clear that this single group was not scalable. So, the concept of API groups was introduced.

Kubernetes uses a versioning scheme to facilitate the evolution of its API. There are three types of versioning in Kubernetes:

Alpha: This is the first stage of the development of a new API. Alpha APIs may be unstable, change significantly after the initial release, and may not even be enabled in your clusters.

Beta: This is the second stage. Beta APIs are well-tested and are enabled by default in your clusters. However, they may still undergo changes, such as in the form of bug fixes or feature enhancements.

Stable: This is the final stage. Stable APIs appear in released software for many subsequent versions.

The version of an API group is represented by `vXalphaX` (e.g., `v1alpha1`), `vXbetaX` (e.g., `v2beta2`), and `vX` (e.g., `v1`) for alpha, beta, and stable versions, respectively.

Kubernetes API groups are divided into two categories

Named API Groups

Named API groups are additional API groups introduced to extend the functionality of Kubernetes beyond the core resources. Each named API group focuses on specific features or functionalities and manages specialized resources related to those features. The Named API group is accessed using the `/apis` endpoint.

apps: This group contains resources related to running applications on Kubernetes. It includes Deployment, ReplicaSet, StatefulSet, and DaemonSet.

batch: This group includes resources for batch processing and job-like tasks. It includes Job and CronJob.

rbac.authorization.k8s.io: This group contains the Role, ClusterRole, RoleBinding, and ClusterRoleBinding resources for handling role-based access control (RBAC) in Kubernetes.

networking.k8s.io: This group contains resources related to networking in k8s, such as NetworkPolicy and Ingress.

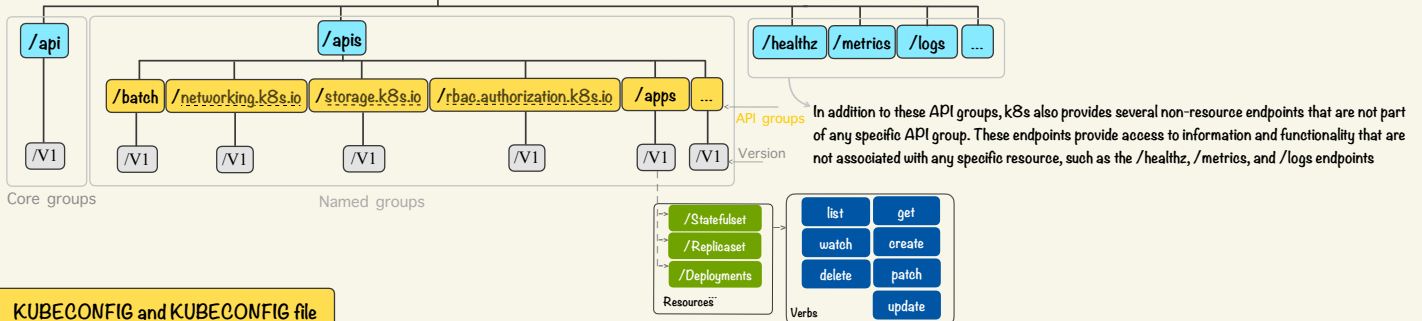
storage.k8s.io: This group contains resources related to storage, such as StorageClass, VolumeAttachment, and the CSI Node driver.

The **core** API group, also referred to as the "v1" group, contains the essential resources that are fundamental to the functioning of a Kubernetes cluster. It includes resources such as Pods, Services, Namespaces, ConfigMaps, Secrets, and more. The core API group is accessed using the `/api/v1` endpoint.

New resources are accessed at the path `/apis/(group)/(version)`. For example, to access the Deployment resource, which is part of the apps group, you would use the path `/apis/apps/v1/deployments`.

Kubernetes APIs

You can list all available API groups and versions in your cluster by running `kubectl api-versions`.



KUBECONFIG and KUBECONFIG file

The `KUBECONFIG` environment variable is used to specify the path to the Kubernetes configuration file, which contains information about the cluster, user, and context used by `kubectl` and other Kubernetes command-line tools. The `KUBECONFIG` file can contain multiple contexts, each representing a different cluster and namespace. The `KUBECONFIG` file is typically stored in the user's home directory at the path `~/.kube/config` on Unix-based systems.

`kubectl --kubeconfig=/path/to/my-kubeconfig/my-kubeconfig.yaml get pods`

This command uses the specified `KUBECONFIG` file instead of the default `~/.kube/config` file.

To create a kubeconfig file using `kubectl`, you can follow these steps:

Set the cluster details **Set the user credentials** **Set the context** Use the context

`kubectl config set-cluster GKE --server=https://k8s-endpoint-6443 --certificate-authority=ca.crt --embed-certs --kubeconfig new.kubeconfig`

Set a cluster entry in kubeconfig Cluster name Server Address Embeds the certificate data directly in the kubeconfig instead of linking to a file kubeconfig file that will be created with this new entry.

`kubectl config set-credentials Arye --client-key=/path/to/arye.key --client-certificate=/path/to/arye.crt --embed-certs --kubeconfig new.kubeconfig`

User name user key file user certificate file If you use a token

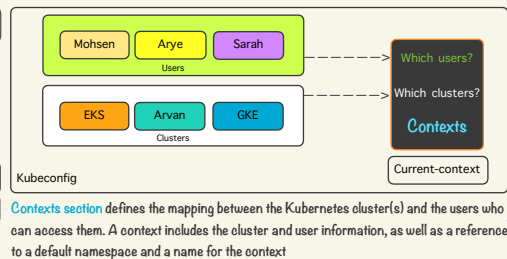
`kubectl config set-credentials Arye --token=eyJhbGciOiJIUzI1NiIsInR5cCI6IkpXLTJr9APOb2rsWHr9HWA --kubeconfig new.kubeconfig`

If you use a client certificate

`kubectl config set-context Arye@GKE --cluster=GKE --user=Arye --namespace=dev --kubeconfig new.kubeconfig`

context name default namespace

`kubectl config use-context Arye@GKE`



Contexts section defines the mapping between the Kubernetes cluster(s) and the users who can access them. A context includes the cluster and user information, as well as a reference to a default namespace and a name for the context.

```
apiVersion: v1
Kind: Config
Clusters:
- name: GKE
  cluster:
    certificate-authority: ca.crt
    server: https://k8s-endpoint-6443
Contexts:
- name: Arye@GKE
  context:
    cluster: GKE
    user: Arye
    namespace: dev
Users:
- name: Arye
  user:
    client-certificate: arye.crt
    client-key: arye.key
current-context: Arye@GKE
```

The default namespace to use for this context.

you can specify a different kubeconfig file by setting the `KUBECONFIG` environment variable.

`export KUBECONFIG=new.kubeconfig`

Cluster Scope in Kubernetes

In Kubernetes, resources are divided into two categories based on their scope: Namespaced and Cluster-scoped.

Namespaced resources: These resources exist and operate within a namespace. They can have different configurations and states in different namespaces.


Cluster-scoped resources: These resources exist and operate across the entire cluster. They are not confined to any particular namespace.

configmaps Role PVC Rolebinding
Deployment ReplicaSet Pods Services Jobs `kubectl api-resources --namespace=true`

nodes PV Clusterroles
Clusterrolebinding Namespace
`kubectl api-resources --namespace=false`

NAME	SHORTNAME	APIVERSION	NAMESPACED	KIND
...	no	v1	true	Endpoints
endpoint	no	v1	false	Pod
Pods	svc	apps/v1	true	Service
deployment	deploy	extensions/v1beta1	true	Deployment
Ingresses	ing		true	Ingress
...				

How to create a new admin or developer user account for accessing to a k8s cluster with X.509 ?



1 Creating a new private key & a csr file by new user

Generate a private key for the user using OpenSSL. The private key is used as part of the user's credentials to authenticate with the Kubernetes API server

```
openssl genrsa -out mojtaba.key 2048
```

Create a CSR for the user using the private key

```
openssl req -new -key mojtaba.key -subj "/CN=mojtaba" -out mojtaba.csr
```

The CSR includes the user's identifying information and the public key associated with the private key

2 Creating a new CSR yaml file and Sign the CSR using the Kubernetes CA

Create a CertificateSigningRequest object in Kubernetes that includes the user's CSR and submit the CSR to the Kubernetes cluster

```

apiVersion: certificates.k8s.io/v1
kind: CertificateSigningRequest
metadata:
  name: mojtaba
spec:
  groups:
  - system:authenticated
  request: LS0hLS1CRUdJTTkKNUedC9BWT...OKLS0hLS1FTkQ...
  signerName: kubernetes.io/kube-apiserver-client
  usages:
  - client auth
  
```

request: \$(cat user-name.csr | base64 -w 0)
Or
cat mojtaba.csr | base64 -w 0

The **signerName** specifies the Kubernetes CA that will sign the certificate.

The **usages** field specifies that the certificate will be used for client authentication.

```
kubectl apply -f csr-mojtaba.yml
```

3 submit the CSR to the Kubernetes cluster and approve it

Once the CSR is submitted, it needs to be approved by a cluster administrator.

```
k get csr
```

NAME	AGE	SIGNERNAME	REQUESTOR	CONDITION
mojtaba	33m	kubernetes.io/kube-apiserver-client	kubernetes-admin	Pending

```
k describe csr mojtaba
```

```

Name: mojtaba
Labels: <none>
Annotations: kubelet.kubernetes.io/last-applied-configuration={"apiVersion":"certificates.k8s.io/v1","kind":"CertificateSigningRequest","metadata":{"name":"mojtaba"},"spec":{"groups":["system:authenticated"],"request":"LS0hLS1CRUdJTTkKNUedC9BWTOKLS0hLS1FTkQgQVVSVEIGSUNBVEUGKVRVUVTVC0hS0tCg==","signerName":"kubernetes.io/kube-apiserver-client","usages":["client auth"]}}
CreationTimestamp: Sat, 11 Jun 2022 17:51:33 +0000
Requesting User: kubernetes-admin
Signer: kubernetes.io/kube-apiserver-client
Status: Pending
Subject: Common Name: mojtaba
Serial Number:
Organization: StarkWare
Organizational Unit: blockchain
Country: IL
Locality: haifa
Province: haifa
Events: <none>
  
```

```
k certificate approve mojtaba
```

```
k get csr
```

NAME	AGE	SIGNERNAME	REQUESTOR	CONDITION
mojtaba	39m	kubernetes.io/kube-apiserver-client	kubernetes-admin	Approved, Issued

This command notifies the Kubernetes CA that the CSR has been approved and requests a signed certificate for the user. The signed certificate is then stored in the `status.certificate` field of the `CertificateSigningRequest` object

4 Export the issued certificate from the CertificateSigningRequest.

you retrieve the signed certificate for the user

```
kubectl get csr mojtaba -o jsonpath='{.status.certificate}' | base64 -d > mojtaba.crt
```

5 Create a kubeconfig File for the User

Create a kubeconfig file for the user that includes the cluster details, user credentials, and context. The `certificate-authority-data` field contains the base64-encoded CA certificate for the Kubernetes cluster.

```

apiVersion: v1
kind: Config

current-context: mojtaba@cka

clusters:
- name: cka
  cluster:
    server: https://kubemaster-1:6443
    certificate-authority: ca.crt

users:
- name: mojtaba
  user:
    client-certificate: mojtaba.crt
    client-key: mojtaba.key

contexts:
- name: mojtaba@cka
  context:
    cluster: cka
    user: mojtaba
    namespace: dev
  
```

```

apiVersion: v1
kind: Config

current-context: mojtaba@cka

clusters:
- name: cka
  cluster:
    server: https://kubemaster-1:6443
    certificate-authority-data: <base64-encoded CA certificate data>

users:
- name: mojtaba
  user:
    client-certificate-data: <base64-encoded client certificate data>
    client-key-data: <base64-encoded client key data>

contexts:
- name: mojtaba@cka
  context:
    cluster: cka
    user: mojtaba
    namespace: dev
  
```

To become independent from external files in the configuration, you can use the data field directly within the configuration file

certificate-authority-data

client-certificate-data

client-key-data

```
cat /etc/kubernetes/pki/ca.crt | base64 -w 0
```

```
cat mojtaba.csr | base64 -w 0
```

```
cat mojtaba.csr | base64 -w 0
```

5.1 If you don't want to create a kubeconfig manually, you can create a kubeconfig using kubectl

```
kubectl config set-cluster cka --server=https://kubemaster-1:6443 --certificate-authority=ca.crt --embed-certs --kubeconfig devuser.kubeconfig
```

```
kubectl config set-credentials mojtaba --client-key=/path/to/mojtaba.key --client-certificate=/path/to/mojtaba.crt --embed-certs --kubeconfig devuser.kubeconfig
```

```
kubectl config set-context mojtaba@cka --cluster=cka --user=mojtaba namespace=dev --kubeconfig devuser.kubeconfig
```

```
kubectl config use-context mojtaba@cka
```

6 Set Up Role-Based Access Control (RBAC) for the User

In this final step, you create a role and role binding to grant the user permissions in the Kubernetes cluster

```

apiVersion: rbac.authorization.k8s.io/v1
kind: Role
metadata:
  name: developer
  namespace: dev
rules:
- apiGroups: [ "" ]
  resources: [ "pods" ]
  verbs: [ "list", "get", "create", "update", "delete" ]
- apiGroups: [ "" ]
  resources: [ "configMap" ]
  verbs: [ "create" ]
  
```

```

apiVersion: rbac.authorization.k8s.io/v1
kind: RoleBinding
metadata:
  name: mojtaba-developer
  namespace: dev
roleRef:
  apiGroup: "rbac.authorization.k8s.io"
  kind: "Role"
  name: "developer"
subjects:
- apiGroup: "rbac.authorization.k8s.io"
  kind: "User"
  name: "mojtaba"
  
```

The reason for having two separate rules in the Role definition is that the two resources, "pods" and "configMap", have different permissions requirements

