


## StatefulSets can use two types of storage

### Shared Storage

All Pods in the StatefulSet share the same storage volume. Data is available to all Pods. Good for things like caches, tmp files etc. Specify a PersistentVolumeClaim template in the sfs spec. All Pods will get a clone of this PVC. data can be corrupted if multiple Pods write to the same files

### Dedicated Storage

Each Pod gets its own PersistentVolume. Data is isolated between Pods. Good for databases, unique files etc. Don't specify volumeClaimTemplates. StatefulSet will create a PVC for each Pod. Updating Pods is harder with dedicated storage, may need to coordinate Pod termination to avoid data loss



```
apiVersion: apps/v1
kind: StatefulSet
metadata:
  name: mysql
spec:
  serviceName: mysql-hs
  replicas: 3
  selector:
    matchLabels:
      app: mysql
```

```
template:
  metadata:
    labels:
      app: mysql
  spec:
    containers:
      - name: mysql
        image: mysql:latest
        env:
          - name: MYSQL_ROOT_PASSWORD
            value: "yourpassword"
        ports:
          - containerPort: 3306
          name: mysql
        volumeMounts:
          - name: mysql-persistent-storage
            mountPath: /var/lib/mysql
```

```
volumes:
  - name: mysql-persistent-storage
    persistentVolumeClaim:
      claimName: mysql-pvc
```

serviceName field specifies the name of the Headless Service that controls the network identity of the StatefulSet's Pods and it is a mandatory field

```
apiVersion: v1
kind: Service
metadata:
  name: mysql-hs
spec:
  ports:
    - port: 3306
  selector:
    app: mysql
  clusterIP: None
```

Pod template

```
apiVersion: v1
kind: PersistentVolumeClaim
metadata:
  name: mysql-pvc
spec:
  accessModes:
    - ReadWriteMany
  storageClassName: manual
  resources:
    requests:
      storage: 1Gi
```

```
apiVersion: v1
kind: PersistentVolume
metadata:
  name: mysql-pv
labels:
  type: local
spec:
  capacity:
    storage: 1Gi
  accessModes:
    - ReadWriteMany
  persistentVolumeReclaimPolicy: Retain
  storageClassName: manual
  nfs:
    path: /srv/nfs/kubedata/pv3
    server: 192.168.49.1
```

The volumeMounts and volumes are defined in the pod template section of the StatefulSet manifest, which means that they will be shared by all pods created by the StatefulSet

The PVC must be created beforehand either manually or through some automated process

The PV and PVC are using the "manual" StorageClass. The PVC has requested a capacity of 1Gi and has been bound to the PV.

k get pv,pvc									
NAME	CAPACITY	ACCESS	MODES	RECLAIM	POLICY	STATUS	CLAIM	STORAGECLASS	AGE
persistentvolume/mysql-pv	1Gi	RWM		Retain		Bound	default/mysql-pvc	manual	5h27m
NAME	STATUS	VOLUME	CAPACITY	ACCESS	MODES	STORAGECLASS	AGE		
persistentvolumeclaim/mysql-pvc-sts3	Bound	mysql-pv	1Gi	RWM		manual	5h29m		

```
apiVersion: apps/v1
kind: StatefulSet
metadata:
  name: mysql-hs
spec:
  selector:
    matchLabels:
      app: mysql
  serviceName: mysql
  replicas: 3
  template:
    metadata:
      labels:
        app: mysql
    spec:
      containers:
        - name: mysql
          image: mysql:latest
          env:
            - name: MYSQL_ROOT_PASSWORD
              valueFrom:
                secretKeyRef:
                  name: mysql-secret
                  key: root-password
          ports:
            - containerPort: 3306
            name: mysql
          volumeMounts:
            - name: mysql-shared-storage
              mountPath: /var/lib/mysql
```

```
volumeClaimTemplates:
  - metadata:
      name: mysql-shared-storage
    spec:
      accessModes: [ "ReadWriteMany" ]
      storageClassName: google-storage
      resources:
        requests:
          storage: 10Gi
```

volumeClaimTemplates is specified at the StatefulSet level, not in the pod template

The "volumeClaimTemplates" field in a StatefulSet is used to define persistent volume claims (PVCs) that will be used by the pods in the set for their storage needs. When a pod is created or rescheduled, it will automatically create/claim one of these PVCs and use it for its persistent storage

If you set the storageClassName to the name of a StorageClass that is configured with a dynamic provisioner, Kubernetes will automatically create a new PV based on the specifications defined in the volumeClaimTemplates section


```
apiVersion: storage.k8s.io/v1
kind: StorageClass
metadata:
  name: google-storage
  provisioner: kubernetes.io/gce-pd
...
```

## How to deploy an application in k8s?

An application in Kubernetes typically consists of YAML files that define the k8s resources needed to run the application, such as Deployments, Services, ConfigMaps, and Secrets. You can deploy the application in Kubernetes manually by creating the YAML files and then using the `kubectl apply` command to create the Kubernetes resources on the cluster. Alternatively, you can use deployment tools like [Kustomize](#), [Helm](#), or the [Helm Operator](#) to automate the deployment process and simplify the creation of the YAML files. These tools provide a higher-level abstraction for managing Kubernetes resources and can make it easier to deploy and manage complex applications in Kubernetes.

### >> Manual Deployment:

Manually deploying applications in Kubernetes involves creating YAML files that define the k8s resources needed to run the application, such as deployments, services, and config maps. You would then use the `kubectl apply` command to create those resources on the Kubernetes cluster.



This approach can be useful for simple applications or for users who prefer a more hands-on approach, but it can be **time-consuming** and **error-prone** for more complex applications.

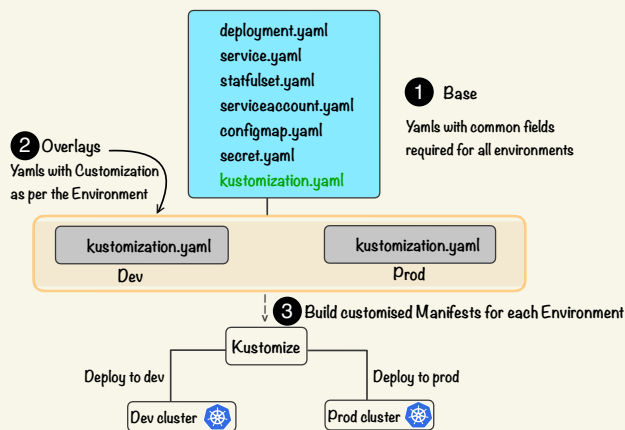
### >> Kustomize

Kustomize is a tool for managing k8s manifest files using a declarative approach. It allows you to define a set of base manifests that define the desired state of your Kubernetes resources, and then apply changes using composition and customization. The basic workflow of Kustomize consists of the following steps:

Create a **base directory** containing your Kubernetes manifests. This directory represents the desired state of your application or environment


Define a **kustomization.yaml** file in the base directory. This file specifies the base resources to use, as well as any additional resources that should be added, modified, or removed

Create overlay **directories** for each **environment** or application variant, if needed. These overlay directories contain additional resources or modifications to apply on top of the base resources



### >> Helm

Helm is a widely-used package manager for Kubernetes that simplifies the deployment and management of applications on a k8s cluster. It enables developers to **package their applications as charts**, which are reusable and shareable bundles that contain all the resources required to deploy an application on a Kubernetes cluster. With Helm, users can easily search for charts, **install and upgrade applications**, **rollback changes**, and **manage dependencies** through a straightforward command-line interface. Additionally, Helm supports versioning, which allows users to track changes to their applications over time and roll back to previous versions if necessary.



Helm uses a packaging format called "charts". A chart is a collection of files that describe a related set of Kubernetes resources. For example, instead of manually creating deployments, services, and other K8s objects, you can package these into a Helm chart. Then, anyone can easily deploy your application by installing the chart.

A Helm chart typically includes the following files:

**Chart.yaml:** This is the core file which includes the name, description, and version of the chart. This file is used by Helm to identify the chart and to provide information to the user when installing or upgrading the chart

```
apiVersion: v2
name: wordpress
description: A Helm chart for deploying WordPress on Kubernetes
version: 1.0.0
appVersion: 5.8.0
maintainers:
- name: Your Name
  email: your@email.com
```

To change a value, you can modify the values.yaml file and then run the helm upgrade command

The template syntax, enclosed in double curly braces ({{}}), is used to reference the values specified in values.yaml

```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: {{ .Release.Name }}-wordpress
  labels:
    app: wordpress
spec:
  replicas: {{ .Values.wordpress.replicaCount }}
  selector:
    matchLabels:
      app: wordpress
  template:
    metadata:
      labels:
        app: wordpress
    spec:
      image: {{ .Values.wordpress.image }}
      imagePullPolicy: {{ .Values.wordpress.imagePullPolicy }}
      ports:
        - name: http
          containerPort: 80
```

**values.yaml:** This file contains the default values for the chart's parameters. These parameters are used in the templates to generate the Kubernetes YAML files. The user can override these values during installation or upgrade using the `--set` flag or a values file. This file is used to allow users to customize the behavior of the chart without modifying the templates directly.

```
wordpress:
  image: wordpress:5.8.0-php7.4-apache
  imagePullPolicy: IfNotPresent
  replicaCount: 1
```

The values in this file are used by the templates in the templates/ directory to generate the k8s YAML files

**templates/:** This directory contains the Kubernetes YAML files that define the resources to be deployed. These files are usually written in a templating language like Go templating or Helm's own template language. The templates can **include placeholders** for the values defined in the values.yaml file. The templates can also include logic to conditionally include or exclude resources based on the values of the parameters

**helpers.tpl:** This file contains reusable snippets of code that can be used in the templates. These snippets can be used to simplify the templates and make them more readable. For example, a helper function might generate a random password or generate a unique name for a resource.